

Anton Bobrov

Software engineer

Location: Moscow

E-mail: baverman@gmail.com

GitHub: <https://github.com/baverman>

Telegram: [@bvrnm](https://t.me/bvrnm)

Profile

I am a software developer since 2002, a Python backend developer since 2007. I like simple implementations and know a couple of tricks how to investigate and fix performance issues. Also, I have a background with a desktop, web frontend, and mobile development.

I prefer straightforward code with minimum abstractions, the clean architecture (strict layers without abstraction leaks), and narrow and deep components (isolated highly cohesive pieces with an encapsulated domain model and minimal public interface).

I'll choose a modular monolith in 99.9% of cases. However, microservices can solve team scaling issues with additional technical complexity.

Some of my projects to look at:

[Fast TSDB backend for Graphite \(https://github.com/baverman/hisser\)](https://github.com/baverman/hisser)

[Python data validation with web in-mind \(https://github.com/baverman/covador\)](https://github.com/baverman/covador)

[Dead simple queue \(https://github.com/baverman/dsq\)](https://github.com/baverman/dsq)

Python:

py2/py3, stdlib, asyncio, Flask/FlaskAdmin, SQLAlchemy, Alembic, Django, aiohttp, requests, Celery, uWSGI, lxml, cryptography, NumPy, Pandas, Cython, CFFI, ctypes, cpu/memory profilers, memory leaks.

DB:

PostgreSQL, Redis, Mongo, SQLite, LMDB, InfluxDB, Clickhouse, Elasticsearch, MySQL, Hive.

Containers:

Docker, LXC, systemd-spawn, KVM, qemu-nbd.

Tools:

Sentry, Grafana, Graphite, Jenkins, Telegraf

Linux administration:

Ansible, Fabric, systemd, nginx, HAProxy, bash scripting, pipelines, CentOS, Ubuntu, Alpine, networking, LVM.

System programming:

Process management, sockets, TCP, UDP, non-blocking IO, low-level protocol parsing.

Clouds:

GCP, AWS, Linode, OpenNebula.

VCS:

Git/StGit, Mercurial, SVN.

Frontend:

HTML, CSS, JavaScript, jQuery, React.

Other languages:

Node, C, Go, Java, Groovy, PHP.

Notable experience

Mar 2018 — Feb 2020 (2 years), CloudLinux, Team Lead, remote

A lead of Imunify360 backend team (3 members). The main service takes care of an event stream from all client installations (thousands of events per second, gigabytes of data per day), pushes it through ETL pipeline to Clickhouse datastore, and handles notifications back to clients.

Stack: python3, asyncio, aiohttp, Redis, Mongo, Clickhouse, HAProxy, TincVPN, Grafana, Sentry, Docker.

Achievements:

- Reworked service architecture to be able to process a growing event stream.
- Implemented app-level real-time monitoring with in/out metrics on each step of ETL path. Statsd, Telegraf, Hisser, Grafana.
- 100% test coverage of whole ETL path, client proto parsing, queue handling, disk storage, ETL processors. As a result, team can deploy new releases directly to production without manual testing on a stage environment.
- Speed up functional tests from 30 to 10 minutes.
- Optimized ansible playbook. Deploy time shortened from 10 to 1 minute.
- Refactored a private admin site for main service (old unmaintainable code with numerous bugs).
- Implemented a fast reverse-ip lookup component.
- 20x speedup of event validation code.
- Reviewed and fix mongo indexes, added partial indexes to fulfill query requirements.
- Tuned HAProxy timeouts and logging, configured upstream checks.
- Moved ETL storage from Hadoop to Clickhouse and implemented Clickhouse automatic table migrations.
- Mesh VPN between production cluster nodes.
- Implemented geo-distributed RBL.

Aug 2016 — Feb 2017 (7 months), Rambler&Co, Team Lead

A lead of Rambler/News team, 6 members, 4 backend and 2 frontend devs. There were 3 services: desktop website + mobile API, mobile site and admin control panel for the content management team. Load: 800 requests per second. My main duties were: solve performance issues and reduce technical debt.

Stack: python2, PostgreSQL, Redis, Django, flask, celery, SQLAlchemy, nginx, uWSGI.

Achievements:

- 2x codebase reduction from 40k to 20k SLOC. Code was in a messy state, endpoints simultaneously used Django/DjangoORM, Django/SQLAlchemy and Flask/SQLAlchemy with duplicated functionality. I've refactored and implemented all endpoints as Flask/SQLAlchemy views.
- Performance optimization. CPU load was reduced by 4 times. The average response latency was reduced by 3 times. I tuned DB indexes, replace complex ORM code (hybrid props and multi-level relation joins which can lead to up to 300 db requests per one http request) with straightforward queries and introduced 2-level caching with uWSGI and Redis.
- Reduced deploy time from 30 to 5 minutes.
- All services were prepared for migration from FreeBSD to Linux.

May 2014 — Jun 2016 (2 years), Zvooq, Team Lead

A lead of the backend team consisting of 6 members. Zvooq is a music streaming service and our team was taking care of mobile/web API, search, editorial team tools, subscriptions, pushes, billing, 3rd party integrations, and backend code for custom ad campaigns.

Stack: python2, FlaskAdmin, PostgreSQL, Elasticsearch, LMDB, Redis, Werkzeug, Celery, SQLAlchemy, nginx, uWSGI, Sentry

Achievements:

- Performance optimization and huge architecture rework. The main service couldn't handle 200rps load before and can process 1000rps after on the same hardware.
- Application real-time monitoring.
- Reduced deploy time from 20 to 1 minute. Implemented CI/CD
- Force unit tests and full coverage. The team can deploy releases after automatic tests directly to production servers.